



Minibatch Training of GNNs with DGL

Quan Gan


How Minibatch Training of Normal Deep Learning Models Work?




Problem in Minibatch Training of a GNN



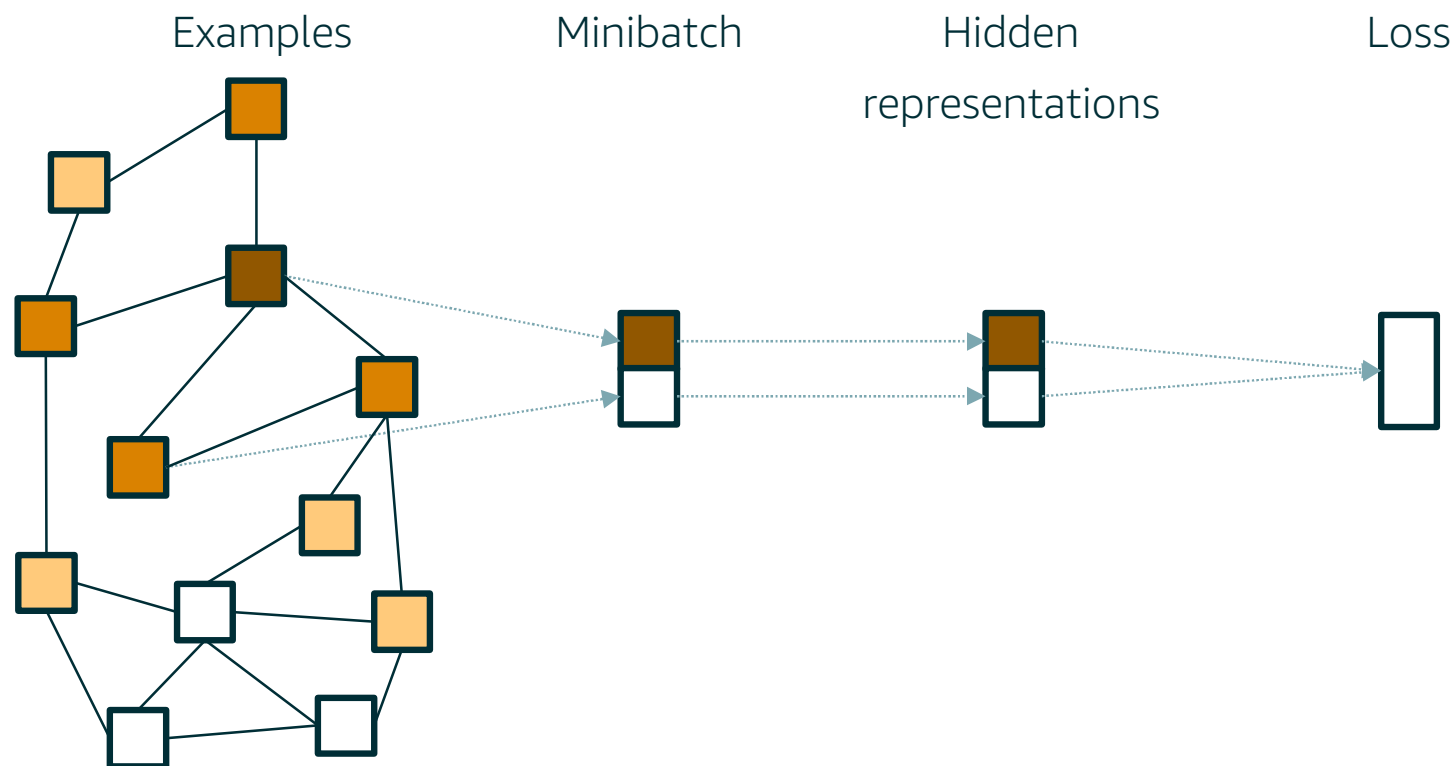
Problem in Minibatch Training of a GNN



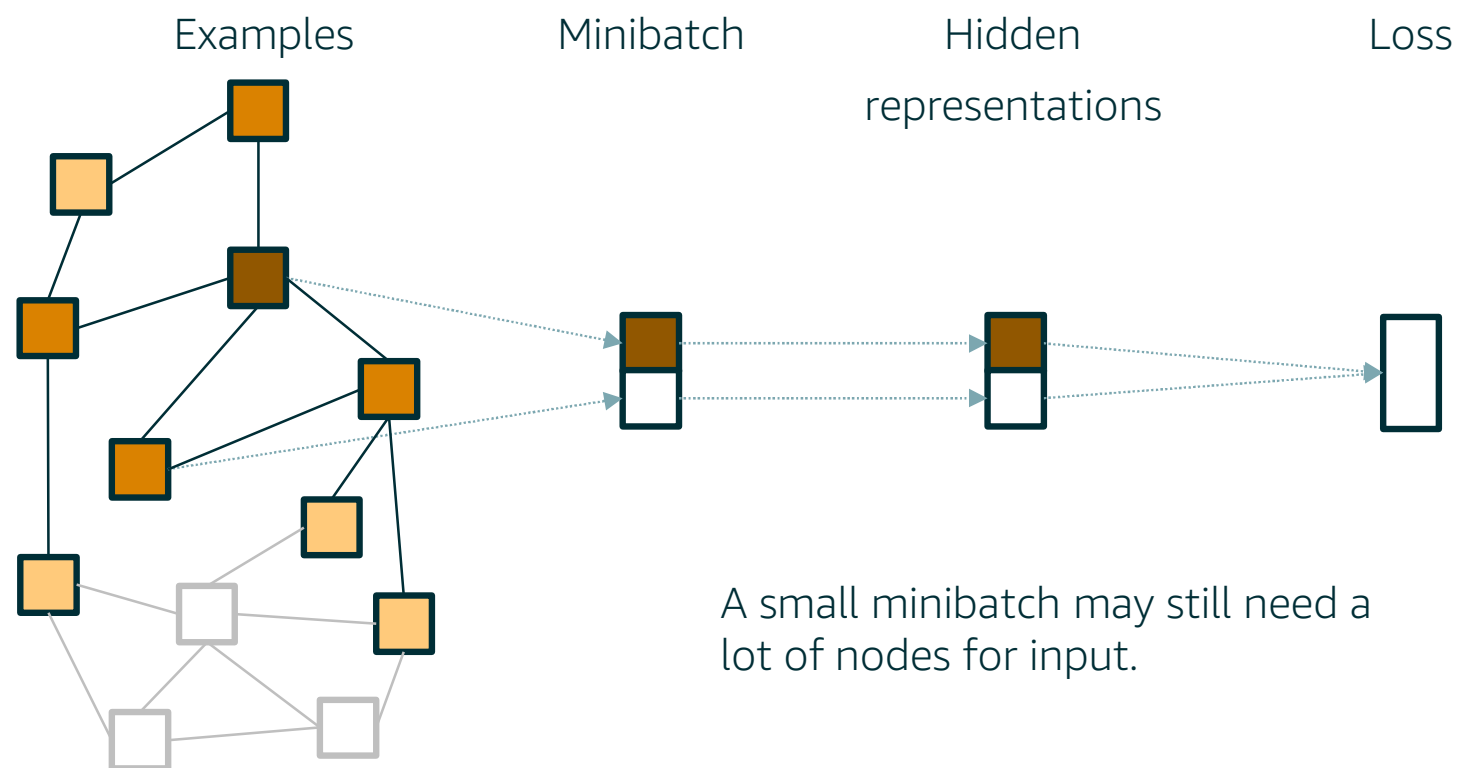
Problem in Minibatch Training of a (1-layer) GNN



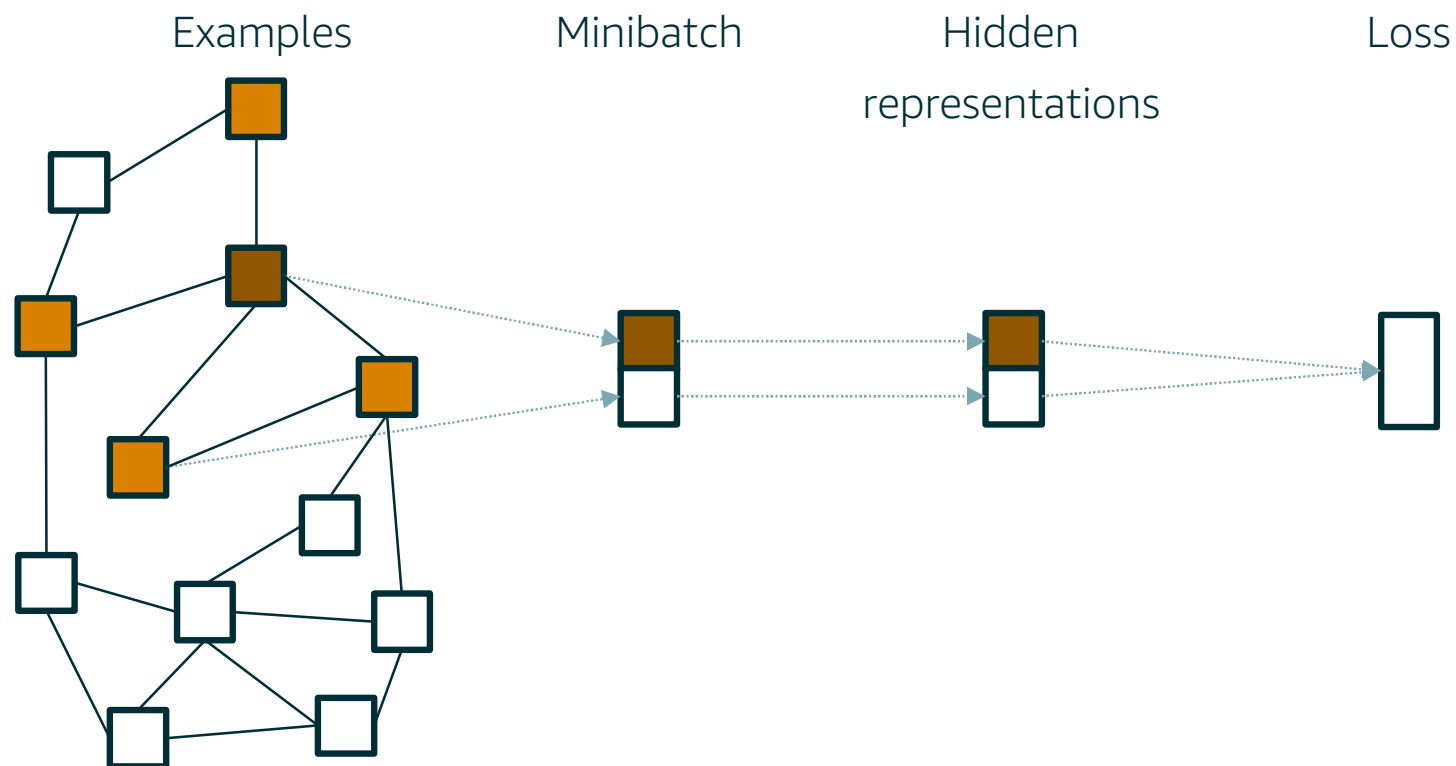
Problem in Minibatch Training of a (2-layer) GNN



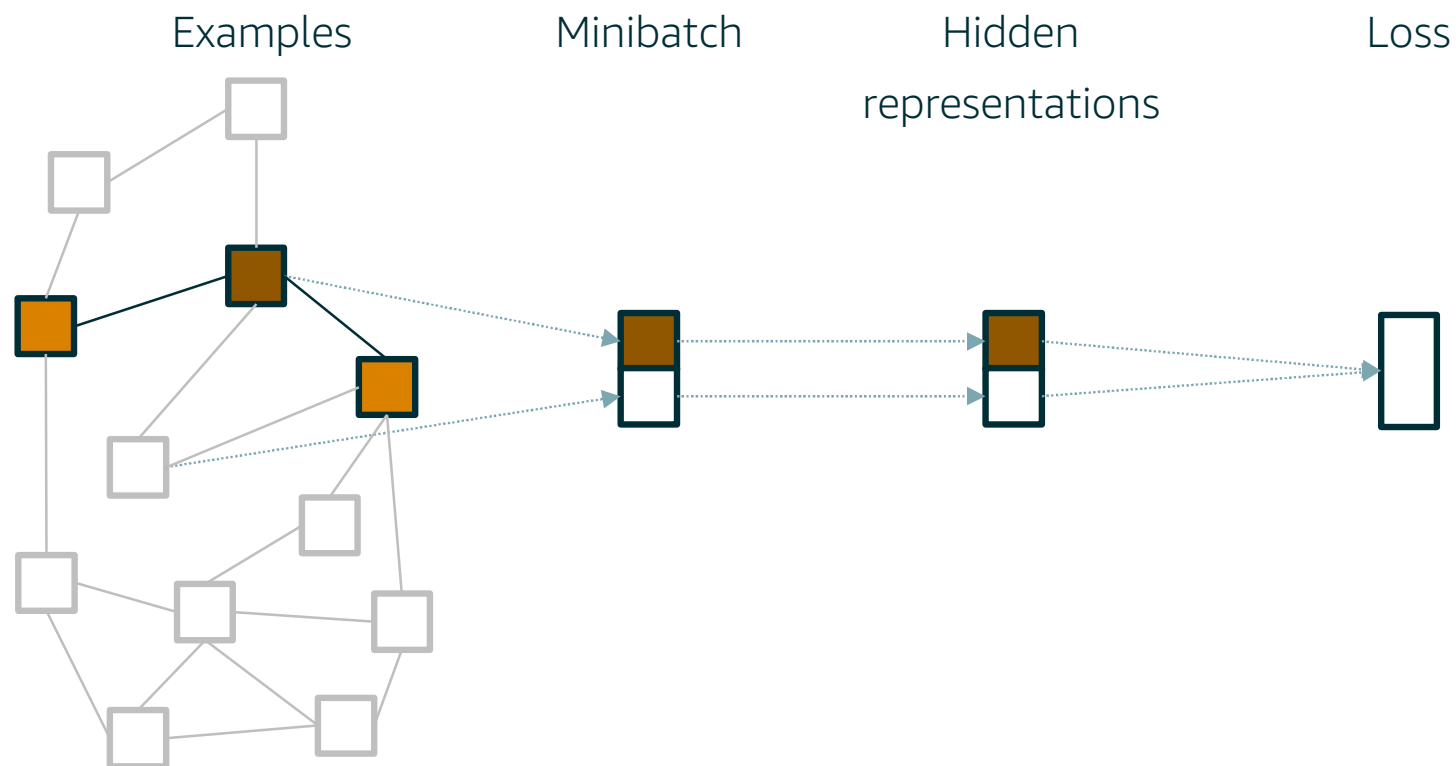
How about using what is necessary only?



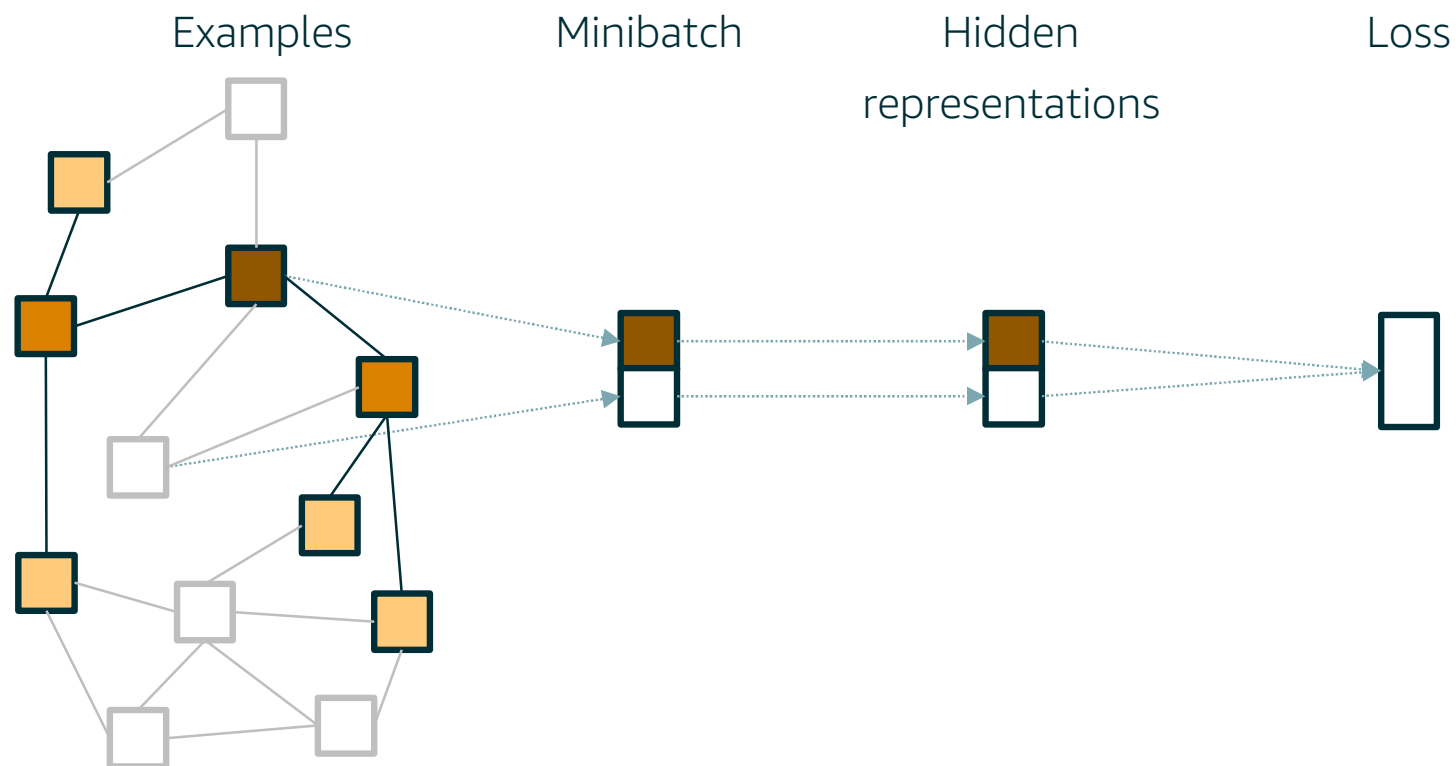
Neighbor Sampling (1-layer)




Neighbor Sampling (1-layer, 2 neighbors per node)




Neighbor Sampling (2-layer, 2 neighbors per node)



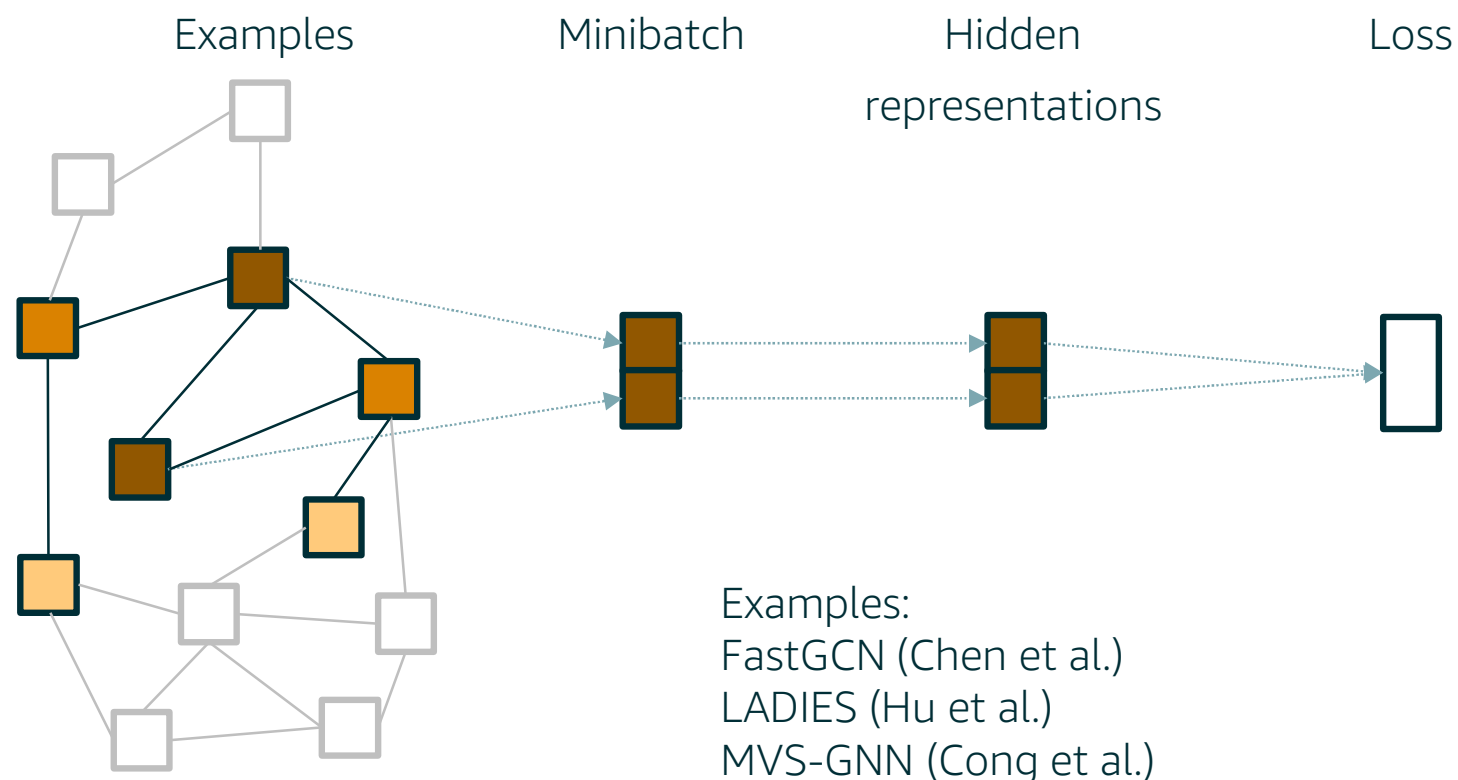
Layerwise Sampling (1-layer)



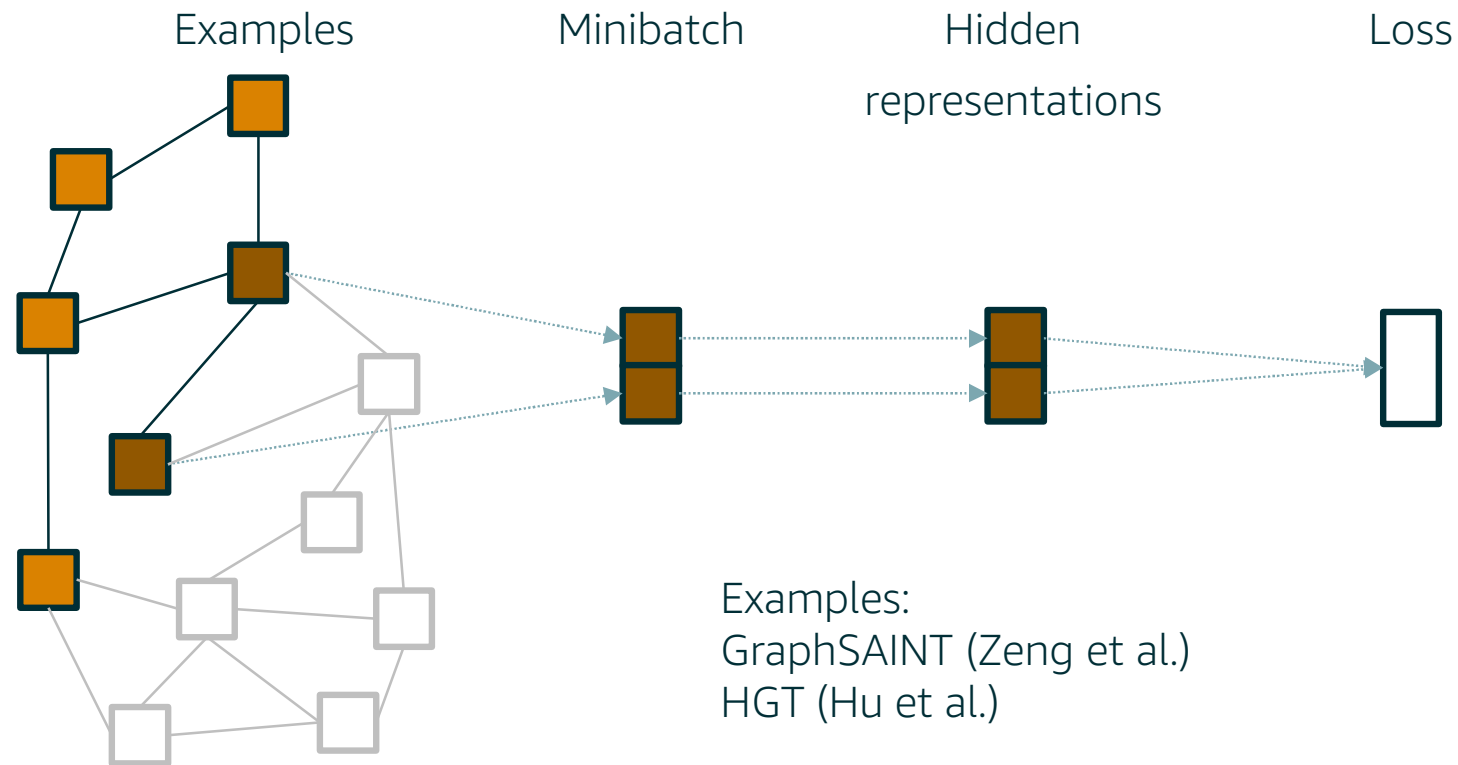
Layerwise Sampling (1-layer, 2 nodes)



Layerwise Sampling (2-layer, 2 nodes per layer)



Subgraph Sampling



Notebook Time!