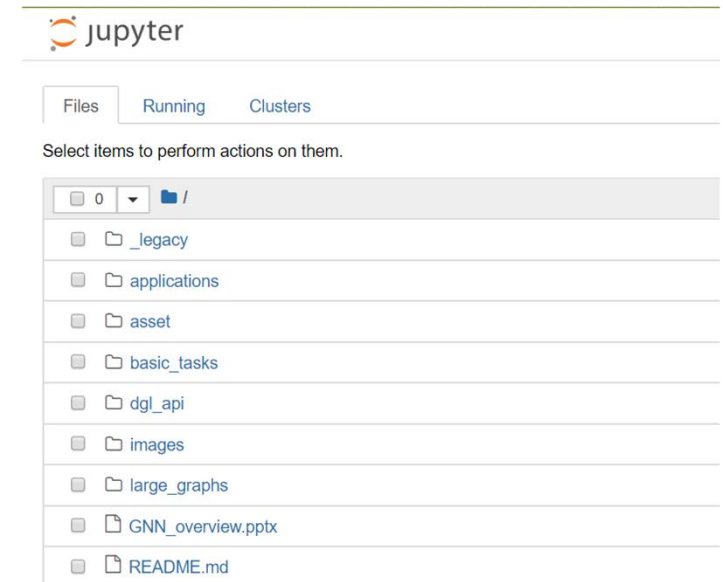
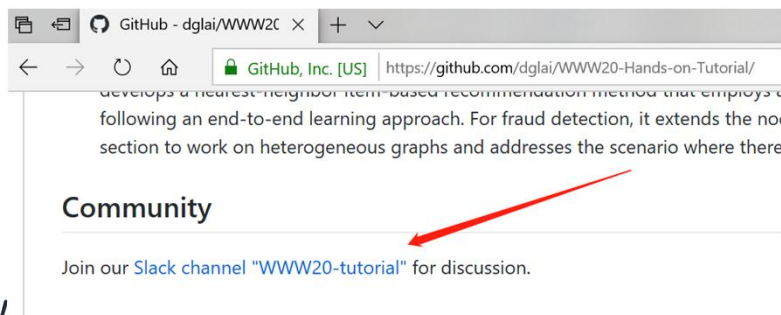


# SageMaker Setup Instructions

- Send an email to [dgl-www@request-nb.mxnet.io](mailto:dgl-www@request-nb.mxnet.io)
- You will receive two emails. Click the notebook link the in the **second** one.

- Need help?
  - Join our slack channel



# Deep Graph Library

## an update

**DGL**



<https://www.dgl.ai>

Zheng Zhang

Director, AWS Shanghai AI Lab ([zhaz@amazon.com](mailto:zhaz@amazon.com))

[zz@nyu.edu](mailto:zz@nyu.edu) (on leave)



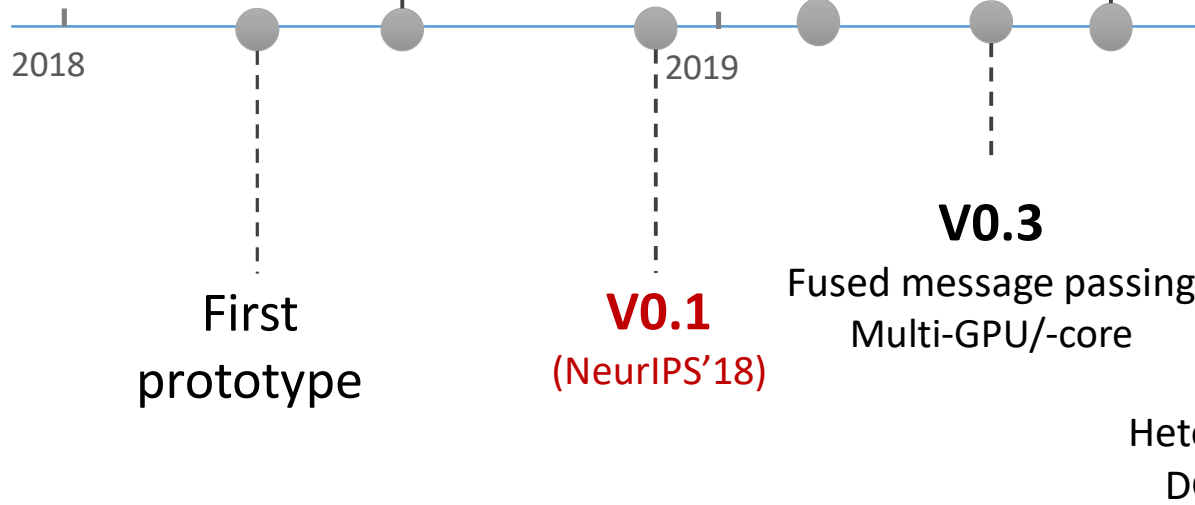
# DGL: The history

Development  
started

**V0.2**  
Sampling APIs

**V0.3.1**  
NN modules  
DGL-Chem

**V0.43**  
TF support  
DGL-KE/DGL-LifeScie  
Sampling



## Introducing Amazon SageMaker Support for Deep Graph Library (DGL): Build and Train Graph Neural Networks

Posted On: Dec 3, 2019

[Amazon SageMaker](#) support for the [Deep Graph Library](#) (DGL) is now available. With DGL, you can improve the prediction accuracy of recommendation, fraud detection, and drug discovery systems using Graph Neural Networks (GNNs).

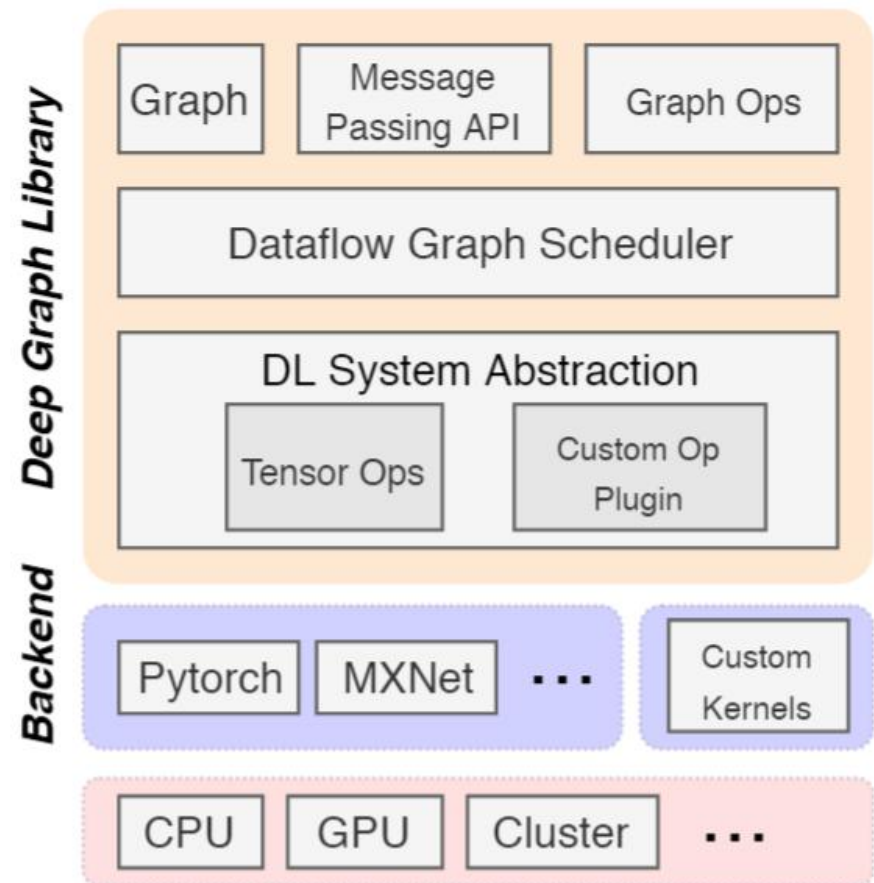
Though GNNs have shown promising results in research, their use in real-world applications has been limited because of the complex infrastructure required to train large graphs and the lack of reliable domain specific models. Developing GNNs involves finding and training on very large graphs with millions of nodes, and it is time-consuming to build and maintain the computational infrastructure required to perform this training. DGL gives you the tools and infrastructure to simplify the implementation and deployment of GNNs.

DGL support in Amazon SageMaker removes the burden of packaging software dependencies, building infrastructure, and finding validated models. As a result, you can test and implement GNNs in hours instead of weeks or months. A Deep Learning container bundles all the software dependencies and the Amazon SageMaker API automatically sets up and scales the infrastructure required to train graphs. With the bundled library of validated models, you can immediately test state-of-the-art GNN models and integrate them into applications.

To get started, please check out the [DGL Get Started](#) page, SageMaker DGL [documentation](#) and our [blog](#).

# DGL meta-objectives & architecture



- Forward and backward compatible
  - **Forward**: easy to develop new models
  - **Backward**: seamless integration with existing frameworks (MXNet/Pytorch/Tensorflow)
- **Fast** and **Scalable**



# DGL: design & API

# Flexible message propagation

- Full propagation (“everyone shouts to everyone near you”)
- Propagation by graph traversal
  - Topological order on sentence parsing tree
  - Belief propagation order
  - Sampling
- Propagation by random walk



# DGL programming interface

- Graph as the core abstraction
  - DGLGraph
  - `g.ndata[ 'h' ]`
- Simple but versatile message passing APIs

$$\text{send}(\mathcal{E}, \phi^e), \quad \text{recv}(\mathcal{V}, \bigoplus, \phi^v)$$

**Active set** specifies which nodes/edges to trigger the computation on.

$\phi^e$   $\phi^v$   $\bigoplus$  can be user-defined functions (**UDFs**) or **built-in** symbolic functions.

# Flexible message handling

*Message function*

$$\text{Edge-wise: } \mathbf{m}_k^{(t)} = \phi^e(\mathbf{e}_k^{(t-1)}, \mathbf{v}_{r_k}^{(t-1)}, \mathbf{v}_{s_k}^{(t-1)}),$$

$$\text{Node-wise: } \mathbf{v}_i^{(t)} = \phi^v(\mathbf{v}_i^{(t-1)}, \bigoplus_{\substack{k \\ \text{s.t. } r_k=i}} \mathbf{m}_k^{(t)}),$$

*Update function*

*Reduce function*

[Gilmer 2017, Wang 2017, Battaglia 2018]



# Writing GNNs is intuitive in DGL

update\_all is a shortcut for  
send(G.edges()) + recv(G.nodes())

```
# code: PyTorch + DGL
# G: DGL Graph
# H: node repr matrix (n_nodes, in_dim)
# W: weights (in_dim * 2, out_dim)
import dgl.function as fn
G.ndata['h'] = H
G.update_all(
    fn.copy_u('h', 'm'),
    fn.max('m', 'h_n'))
H_N = G.ndata['h_n']
H = torch.relu(torch.cat([H_N, H], 1) @ W)
```

$$h_v^{(t+1)} = \max_{u \in \mathcal{N}(v)} h_u^{(t)}$$

```
# code: PyTorch + DGL
# G: DGL Graph
) # H: node repr matrix (n_nodes, in_dim)
# W: weights (in_dim * 2, out_dim)
import dgl.function as fn
G.ndata['h'] = H
G.update_all(
    fn.copy_u('h', 'm'),
    fn.mean('m', 'h_n'))
H_N = G.ndata['h_n']
H = torch.relu(torch.cat([H_N, H], 1) @ W)
```

$$h_v^{(t+1)} = \frac{1}{|\mathcal{N}(v)|} \sum_{u \in \mathcal{N}(v)} h_u^{(t)}$$

# Writing GNNs is intuitive in DGL (GAT)

```
def msg_func(edges):
    h_src = edges.src['h']
    h_dst = edges.dst['h']
    alpha_hat = MLP(
        torch.cat([h_dst, h_src], 1))
    return {'m': h_src, 'alpha_hat': alpha}


def reduce_func(nodes):
    # Incoming messages are batched along
    # 2nd axis.
    m = nodes.mailbox['m']
    alpha_hat = nodes.mailbox['alpha_hat']
    alpha = torch.softmax(alpha_hat, 1)
    return {'h_n':
        (m * alpha[:, None]).sum(1)}
```

```
# code: PyTorch + DGL
# G: DGL Graph
# H: node repr matrix (n_nodes, in_dim)
# W: weights (in_dim * 2, out_dim)

import dgl.function as fn
G.ndata['h'] = H
G.update_all(msg_func, reduce_func)
H_N = G.ndata['h_n']
H = torch.relu(torch.cat([H_N, H], 1) @ W)
```

$$\alpha_{ij} = \frac{\exp\left(\text{LeakyReLU}\left(\vec{a}^T[\mathbf{W}\vec{h}_i \parallel \mathbf{W}\vec{h}_j]\right)\right)}{\sum_{k \in \mathcal{N}_i} \exp\left(\text{LeakyReLU}\left(\vec{a}^T[\mathbf{W}\vec{h}_i \parallel \mathbf{W}\vec{h}_k]\right)\right)}$$

# Different scenarios require different supports




Sampling

Single giant graph



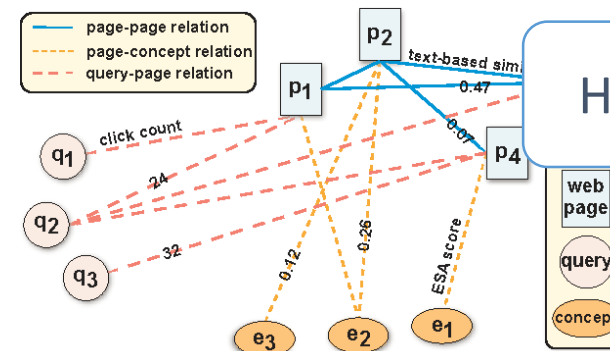
Batching  
graphs

Many moderate-sized graphs



Mutation

Dynamic graph



Heterogeneous

# Summary of DGL Programming Interface

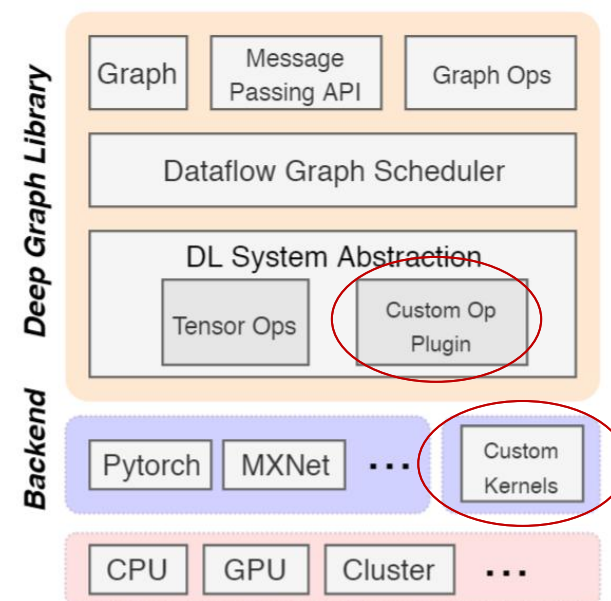
	DGL			
Message Passing	arbitrary $\phi^e$	✓	User-defined functions	Send, Recv
	arbitrary $\phi^v$	✓		
	arbitrary $\oplus$	✓		
Propagation Order	full	✓	Active set	
	partial	✓		
	random walk	✓		
	sampling	✓		
Graph Type	many & small	✓	Batching APIs	Mutation APIs
	single & giant	✓		
	dynamic	✓	Sampling APIs	Support multi-type
	heterogeneous	✓		
System	multi-platform	✓		

# Performance evaluation

# Initially paid a price for multi-backend (woops!)

Table 4: Training runtime comparison.

Dataset	Method	DGL DB	DGL GS	PyG
Cora	GCN	4.19s	0.32s	<b>0.25s</b>
	GAT	6.31s	5.36s	<b>0.80s</b>
CiteSeer	GCN	3.78s	0.34s	<b>0.30s</b>
	GAT	5.61s	4.91s	<b>0.88s</b>
PubMed	GCN	12.91s	0.36s	<b>0.32s</b>
	GAT	18.69s	13.76s	<b>2.42s</b>
MUTAG	RGCN	18.81s	2.40s	<b>2.14s</b>



[Fast Graph Representation Learning with PyTorch Geometric](#)

# Evaluation: efficiency and memory consumption


Dataset	Model	Accuracy	Time		Memory	
			PyG	DGL	PyG	DGL
Cora	GCN	$81.31 \pm 0.88$	<b>0.478</b>	0.666	1.1	1.1
	GAT	$83.98 \pm 0.52$	1.608	<b>1.399</b>	1.2	<b>1.1</b>
CiteSeer	GCN	$70.98 \pm 0.68$	<b>0.490</b>	0.674	1.1	1.1
	GAT	$69.96 \pm 0.53$	1.606	<b>1.399</b>	1.3	<b>1.2</b>
PubMed	GCN	$79.00 \pm 0.41$	<b>0.491</b>	0.690	1.1	1.1
	GAT	$77.65 \pm 0.32$	1.946	<b>1.393</b>	1.6	<b>1.2</b>
Reddit	GCN	$93.46 \pm 0.06$	<i>OOM</i>	<b>28.6</b>	<i>OOM</i>	<b>11.7</b>
Reddit-S	GCN	N/A	29.12	<b>9.44</b>	15.7	<b>3.6</b>

Table 2: Training time (in seconds) for 200 epochs and memory consumption (GB).

Testbed: one V100 GPU (16GB)

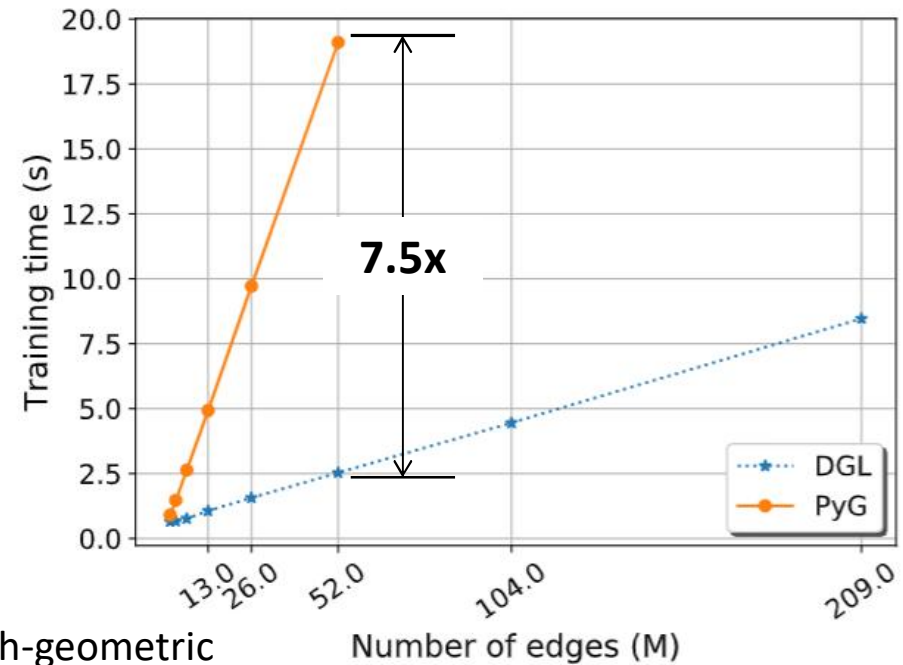


# Scalability: single machine, single GPU



PyG: pytorch-geometric

Scalability with graph size




Scalability with graph density




# Scalability: single machine, multi-GPU

Training speed (K samples/sec)



p3.16xlarge, 8 V100 GPUs, 64 vCPU  
Data set: Reddit (232K nodes, 114M edges)  
GraphSage neighbor sampling

# Scalability: single machine, NUMA




X1, 2TB, 128 vCPU

Data set: Reddit (232K nodes, 114M edges)

Controlled-variate sampling

# Evaluation: auto-batching



Compare with DyNet for training TreeLSTM

Testbed: one V100 GPU (16GB)

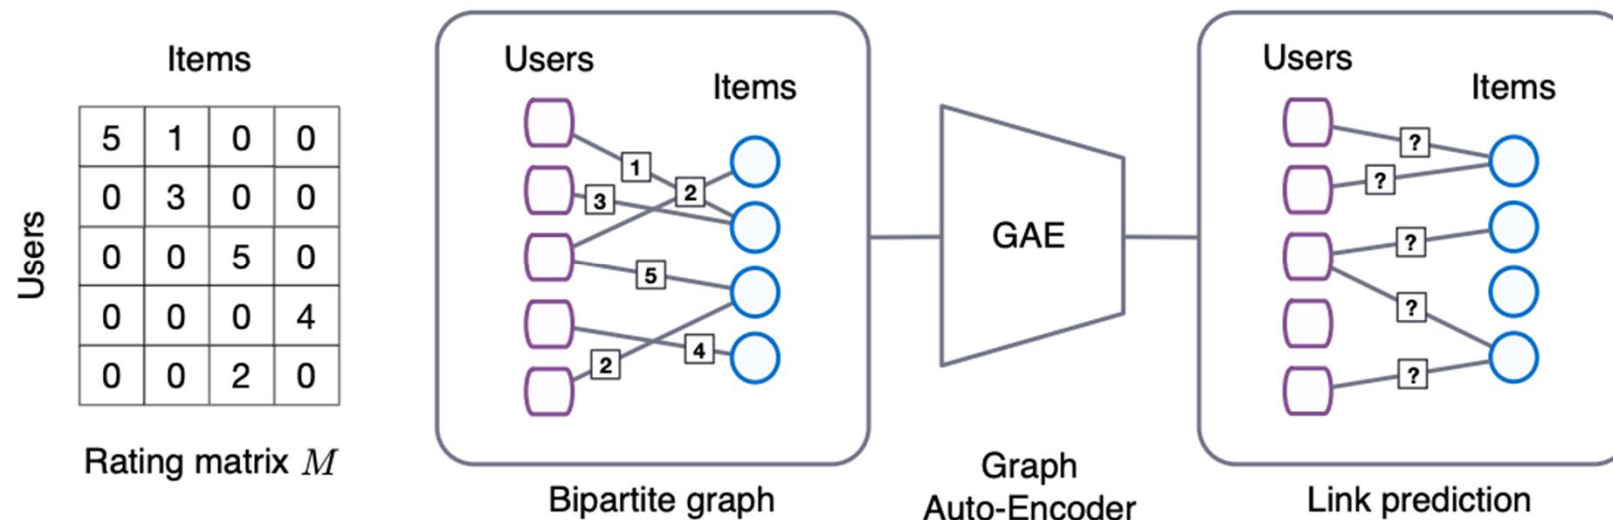
# What's new

Release 0.4 and 0.43



# Heterogenous graph

- Example: recommendation system, GCMC



Graph Convolutional Matrix Completion

# Supporting Heterogeneous Graph

```

1 import torch
2 import torch.nn as nn
3 import torch.nn.functional as F
4 import dgl.function as fn
5
6 class HeteroRGCNLayer(nn.Module):
7     def __init__(self, in_size, out_size, etypes):
8         super(HeteroRGCNLayer, self).__init__()
9         # define parameter W_r for each relation
10        self.weight = nn.ModuleDict({
11            name : nn.Linear(in_size, out_size) for name in etypes
12        })
13
14    def forward(self, G, feat_dict):
15        # G is a heterogeneous graph
16        # feat_dict is a dictionary of features of each node type
17        funcs = {}
18        for srctype, etype, dsttype in G.canonical_etypes:
19            # Compute W_r * h
20            Wh = self.weight[etype](feat_dict[srctype])
21            # Save it to graph
22            G.nodes[srctype].data['Wh_%s' % etype] = Wh
23            # Per-type message passing: (message_func, reduce_func)
24            # All reducers write to the same field 'h', which is a hint for type-wise reducer.
25            funcs[etype] = (fn.copy_u('Wh_%s' % etype, 'm'), fn.mean('m', 'h'))
26        # Trigger message passing on heterograph using multi_update_all
27        # Argument#1: per-type message passing functions.
28        # Argument#2: type-wise reducer, could be: "sum", "max", "min", "mean", "stack"
29        G.multi_update_all(funcs, 'sum')
30        # Return the updated features of each node type.
31        return {ntype : G.nodes[ntype].data['h'] for ntype in G.ntypes}

```

$$h_i^{(l+1)} = \sigma \left( \sum_{r \in \mathcal{R}} \sum_{j \in \mathcal{N}_r(i)} W_r^{(l)} h_j^{(l)} \right)$$

# Example: graph convolutional matrix completion

Dataset	RMSE (DGL)	RMSE (Official)	Speed (DGL)	Speed (Official)	Speedup
MovieLens-100K	<b>0.9077</b>	0.910	<b>0.0246</b> s/epoch	0.1008 s/epoch	<b>5x</b>
MovieLens-1M	0.8377	<b>0.832</b>	<b>0.0695</b> s/epoch	1.538 s/epoch	<b>22x</b>
MovieLens-10M	0.7875	<b>0.777*</b>	<b>0.6480</b> s/epoch	Long*	

\*Official training on MovieLens-10M has to be in mini-batch, which lasts for over 24+ hours

# Example: drug discovery

Training time per epoch (1 V100 GPU)

	DGL	Official	Speedup
<u>ACNN</u>	0.36	1.58	4.4x

[Atomic Convolutional Neural Network \(ACNN\)](#)



# Digress: even more examples of drug discovery

Training time per epoch (1 V100 GPU)

	DGL	Official	Speedup
<a href="#">ACNN</a>	0.36	1.58	4.4x
<a href="#">JTNN</a>	743	1826	2.5x
<a href="#">GCN (Tox21)</a>	1.9	8.4	4.4x
<a href="#">AttentiveFP*</a>	1.2	6.0	5x

\*[Pushing the Boundaries of Molecular Representation for Drug Discovery with the Graph Attention Mechanism](#) (GAT, deep, GRU aggregation)

# DGL 0.43 Release

Minor release loaded with features:

- TF backend support
- 2 DGL domain packages
  - DGL-KGE
  - DGL-LifeSci (5 new models)
- Sampling API

## 0.43 – (1) TensorFlow backend support


- 15 common GNN modules (including heterography support)
- On average x3.5 faster than GraphNet, x1.9 than tf-geometric

Dateset	Model	DGL	GraphNet	tf_geometric
Cora	GCN	<b>0.0148</b>	0.0152	0.0192
Reddit	GCN	<b>0.1095</b>	OOM	OOM
PubMed	GCN	<b>0.0156</b>	0.0553	0.0185
PPI	GCN	<b>0.09</b>	0.16	0.21
Cora	GAT	<b>0.0442</b>	n/a	0.058
PPI	GAT	<b>0.398</b>	n/a	0.752

## 0.43 – (2) DGL-KE:

**Light-speed** for knowledge graph embedding learning, at **scale**

- 2~5 speedup over PyTorch Big Graph/GraphVite




**Full FreeBase**

- **86M** nodes
- **348M** edges
- **30min** on 4 48core machines

## 0.43 – (2) DGL-KE: runs on multi-GPU, too

**Light-speed** for knowledge graph embedding learning, with **scale**

- 2~5 speedup over PyTorch Big Graph/GraphVite



**Full FreeBase**

- **86M** nodes
- **348M** edges
- **100min** on 8-GPU

## 0.43 – (3) Sampling (for giant graph)

```
class NeighborSampler(object):
    def __init__(self, g, fanouts):
        self.g = g # The full graph structure
        self.fanouts = fanouts # fan-out of each layer


    def sample_blocks(self, seeds):
        # `seeds` are the set of nodes to build one sample from.
        blocks = []
        for fanout in self.fanouts:
            # For each seed node, sample ``fanout`` neighbors.
            frontier = dgl.sampling.sample_neighbors(g, seeds, fanout, replace=True)
            # Then we compact the frontier into a bipartite graph for message passing.
            block = dgl.to_block(frontier, seeds)
            # Obtain the seed nodes for next layer.
            seeds = block.srcdata[dgl.NID]

            blocks.insert(0, block)
        return blocks
```

Neighbor sampling, +control variate, random walk ...

- Works with **many GNNs** (PinSAGE, GraphSAGE, GCMC, ...)
- Support **customization** in Python.
- Support **heterogeneous** graphs.
- Use **existing NN modules** without code change.
- **Multi-processing & multi-threading** for maximum speed.

# Distributed training: GCN (preliminary results)



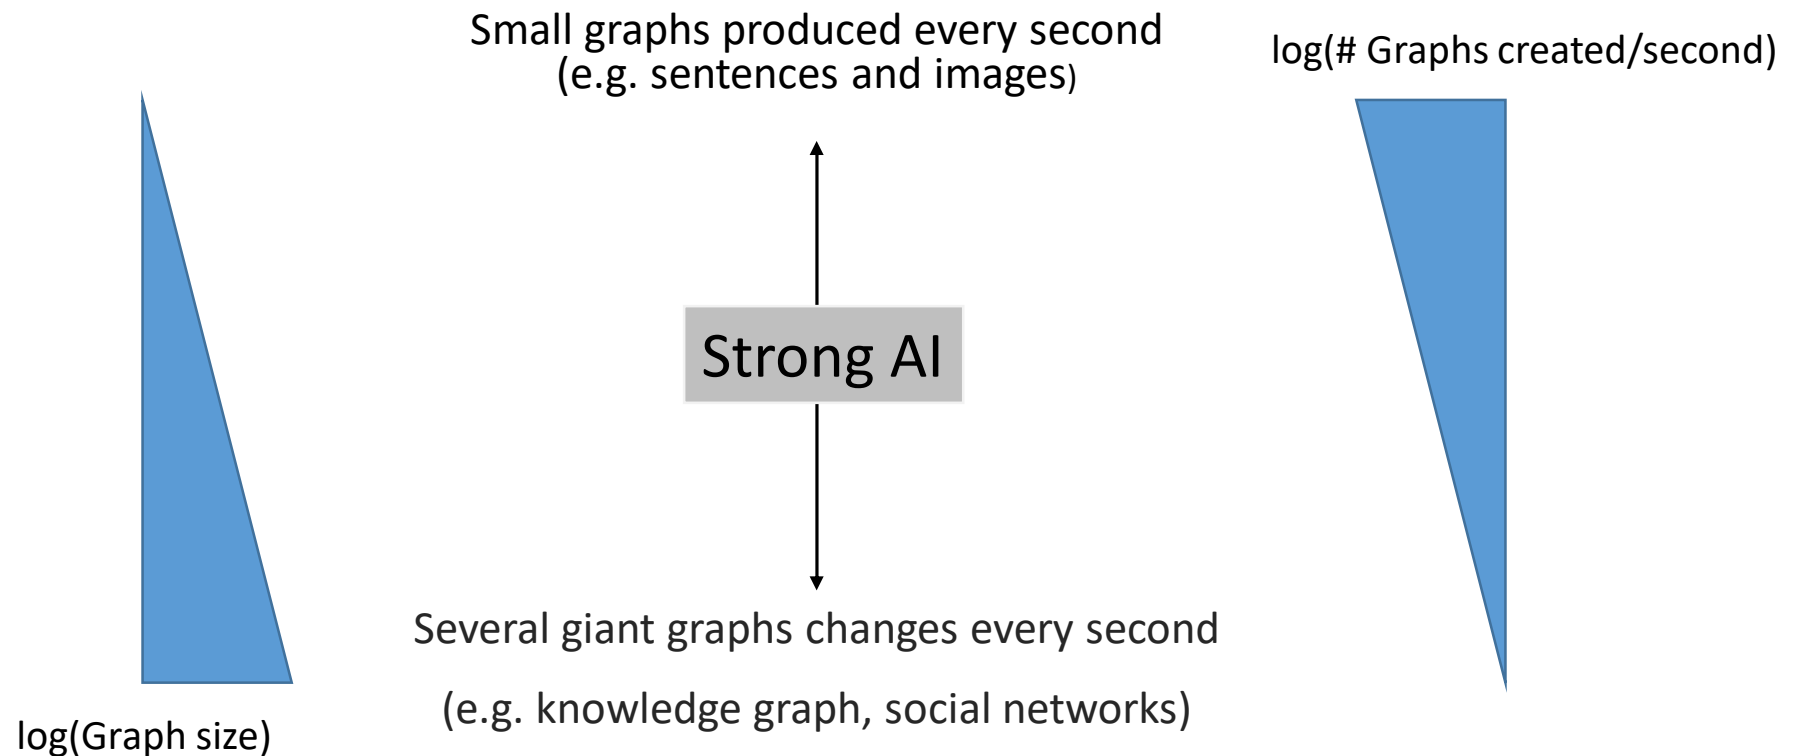
Distributed training of GCN on Reddit dataset.

Neighbor sampling

Data set: Reddit (232K nodes, 114M edges)

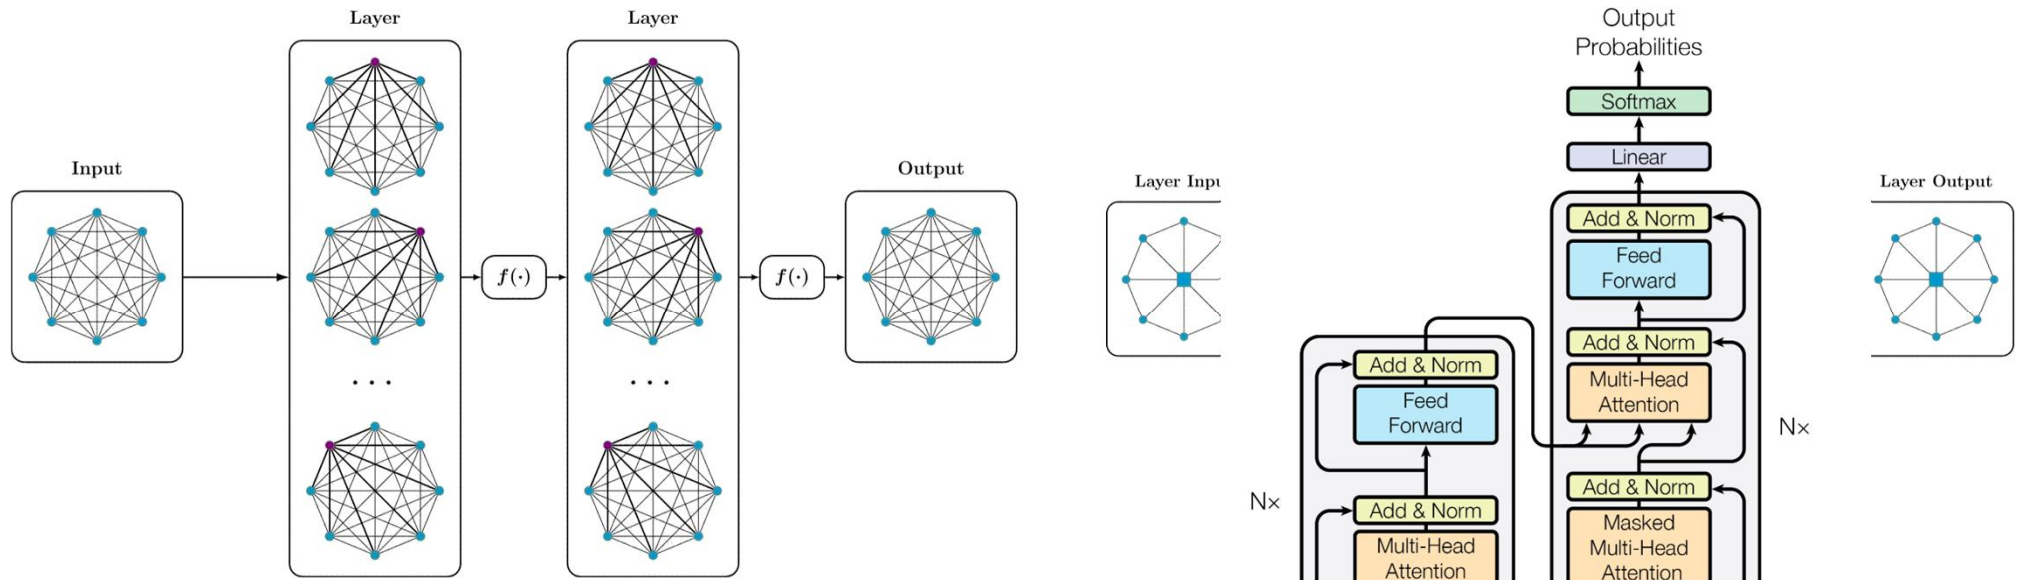
Testbed: c5n.18x, 100Gb/s network, 72vCPU

# Seeing the world from the lens of graphs





# Transformer is GAT (over a complete graph)



Transformer:  
-  $O(n^2)$ : data hungry

Star-Transform  
-  $O(n)$ : mu  
- Leverages

SegTree-Transformer (in ICLR'19 RLGM)

- $O(n \log n)$ : less data hungry
- A good compromise in between

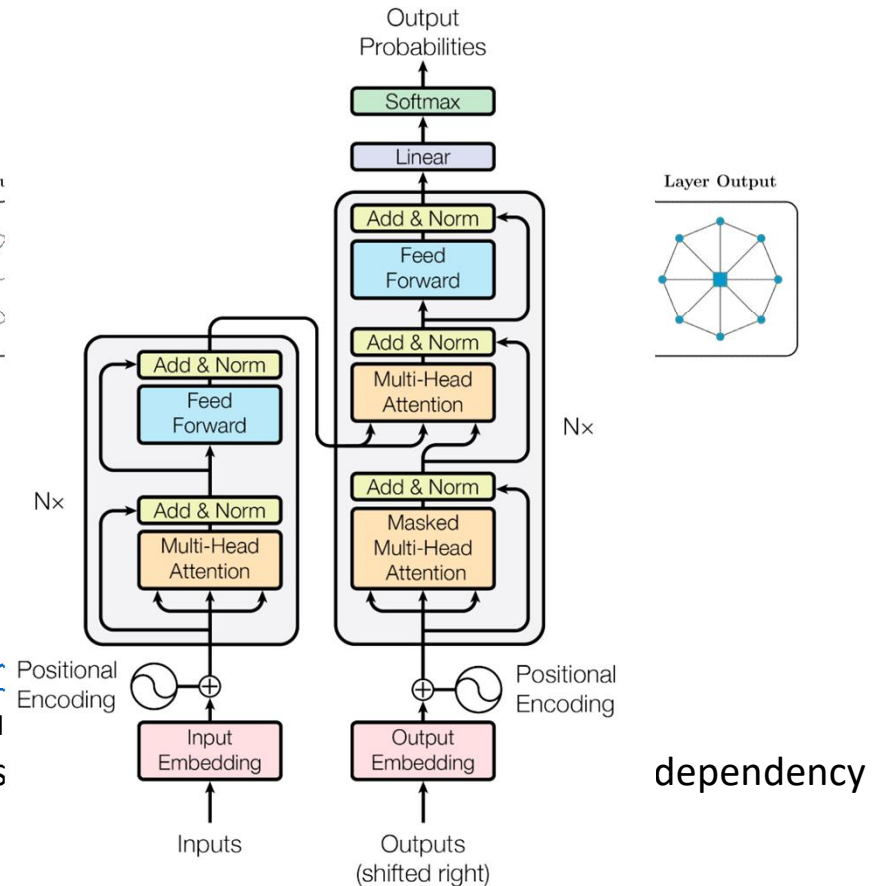
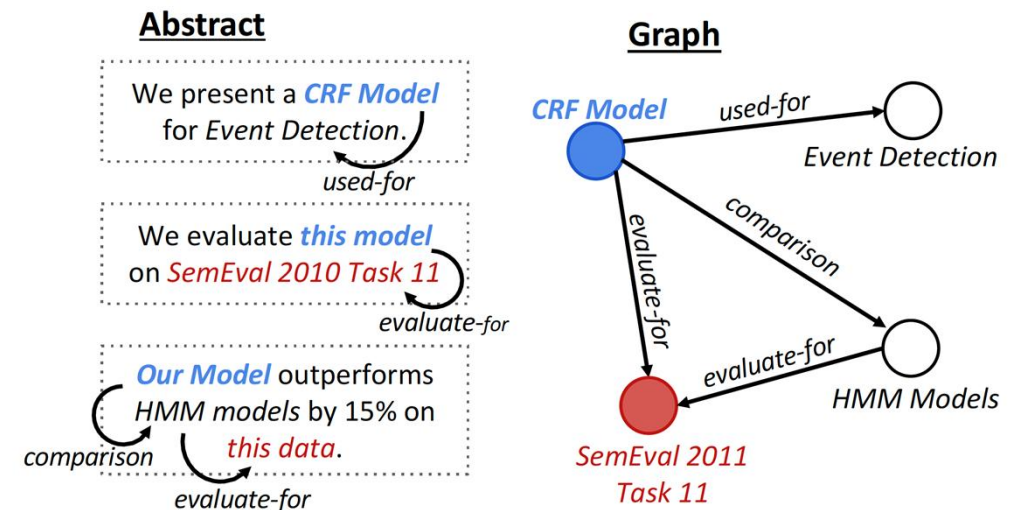


Figure 1: The Transformer - model architecture.

# Graph in Natural Language Processing

- Knowledge Graph to Sentence
  - Input: a title and knowledge graph
  - Output: abstract text

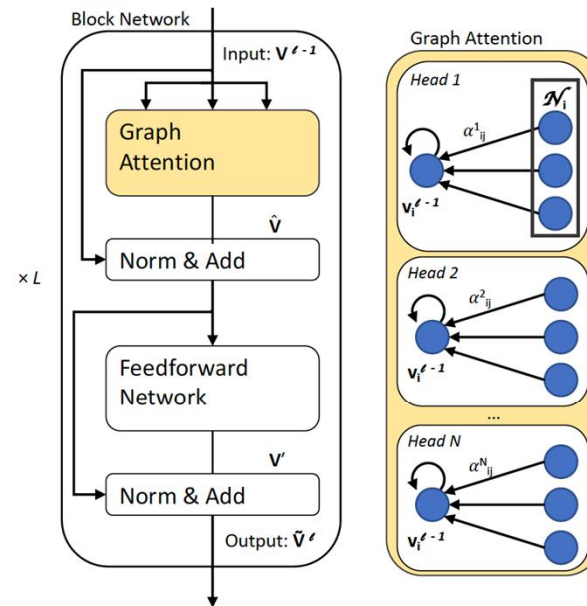
Title: Event Detection with Conditional Random Fields



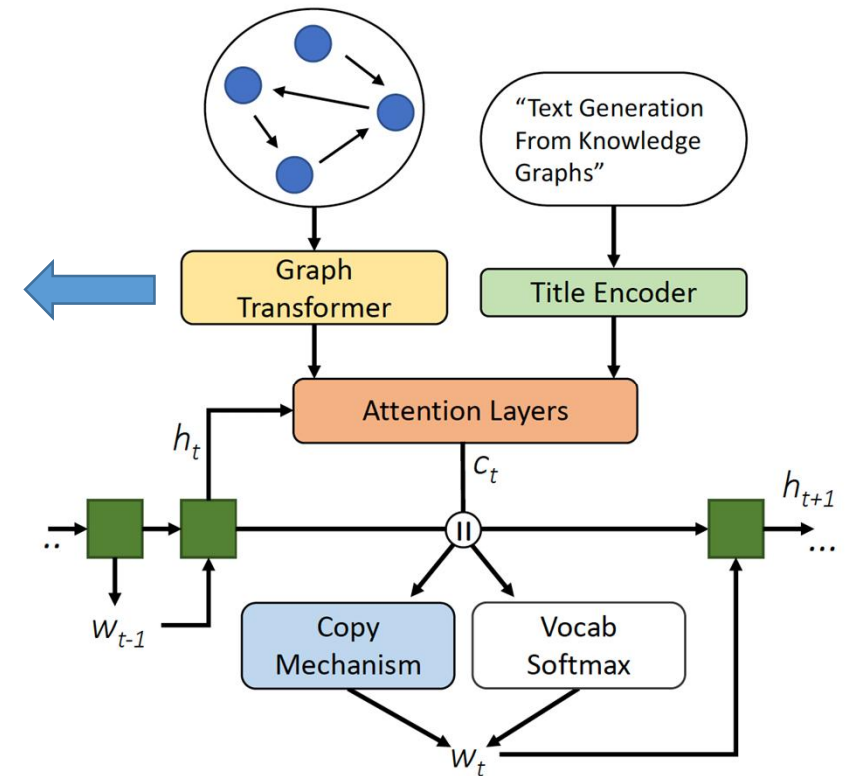
Text Generation from Knowledge Graphs with Graph Transformers

# Graph in Natural Language Processing

- GraphWriter
  - Graph Attention
  - Copy or vocab




DGL	Official	Speedup
1192 (s)	1970 (s)	1.7x




Text Generation from Knowledge Graphs with Graph Transformers

# Scene graph extraction (experimental)

- [Graph R-CNN for Scene Graph Generation](#)
  - Per epoch time 45min (vs. 2.5hr)



Object Detection



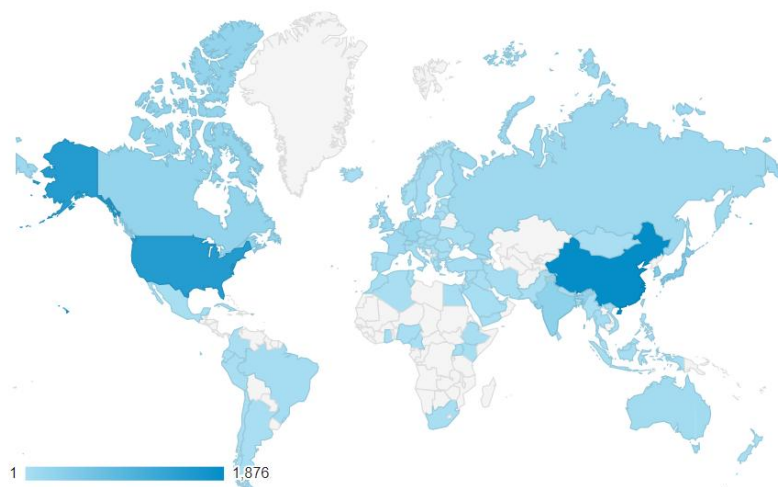

Link exists with prob 68.99%

Link type:

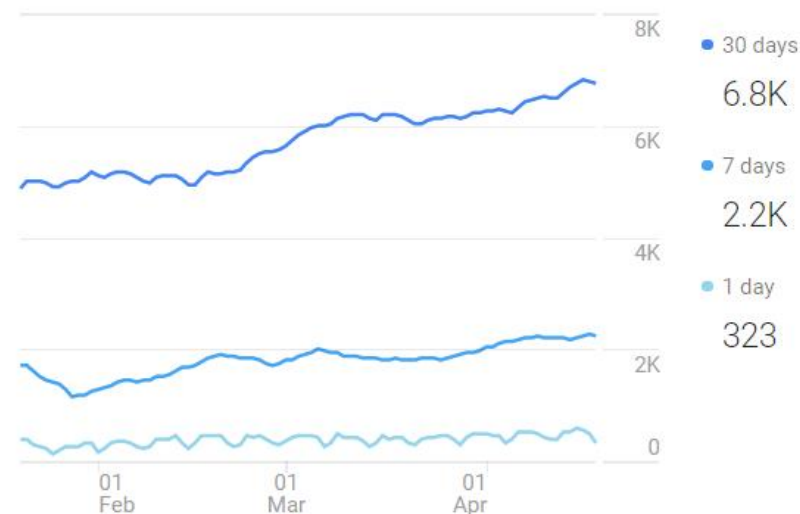
<man> <have> <shirt>, 43.29%

<man> <wear> <shirt>, 39.64%

# Open source, the source of innovation



Active Users

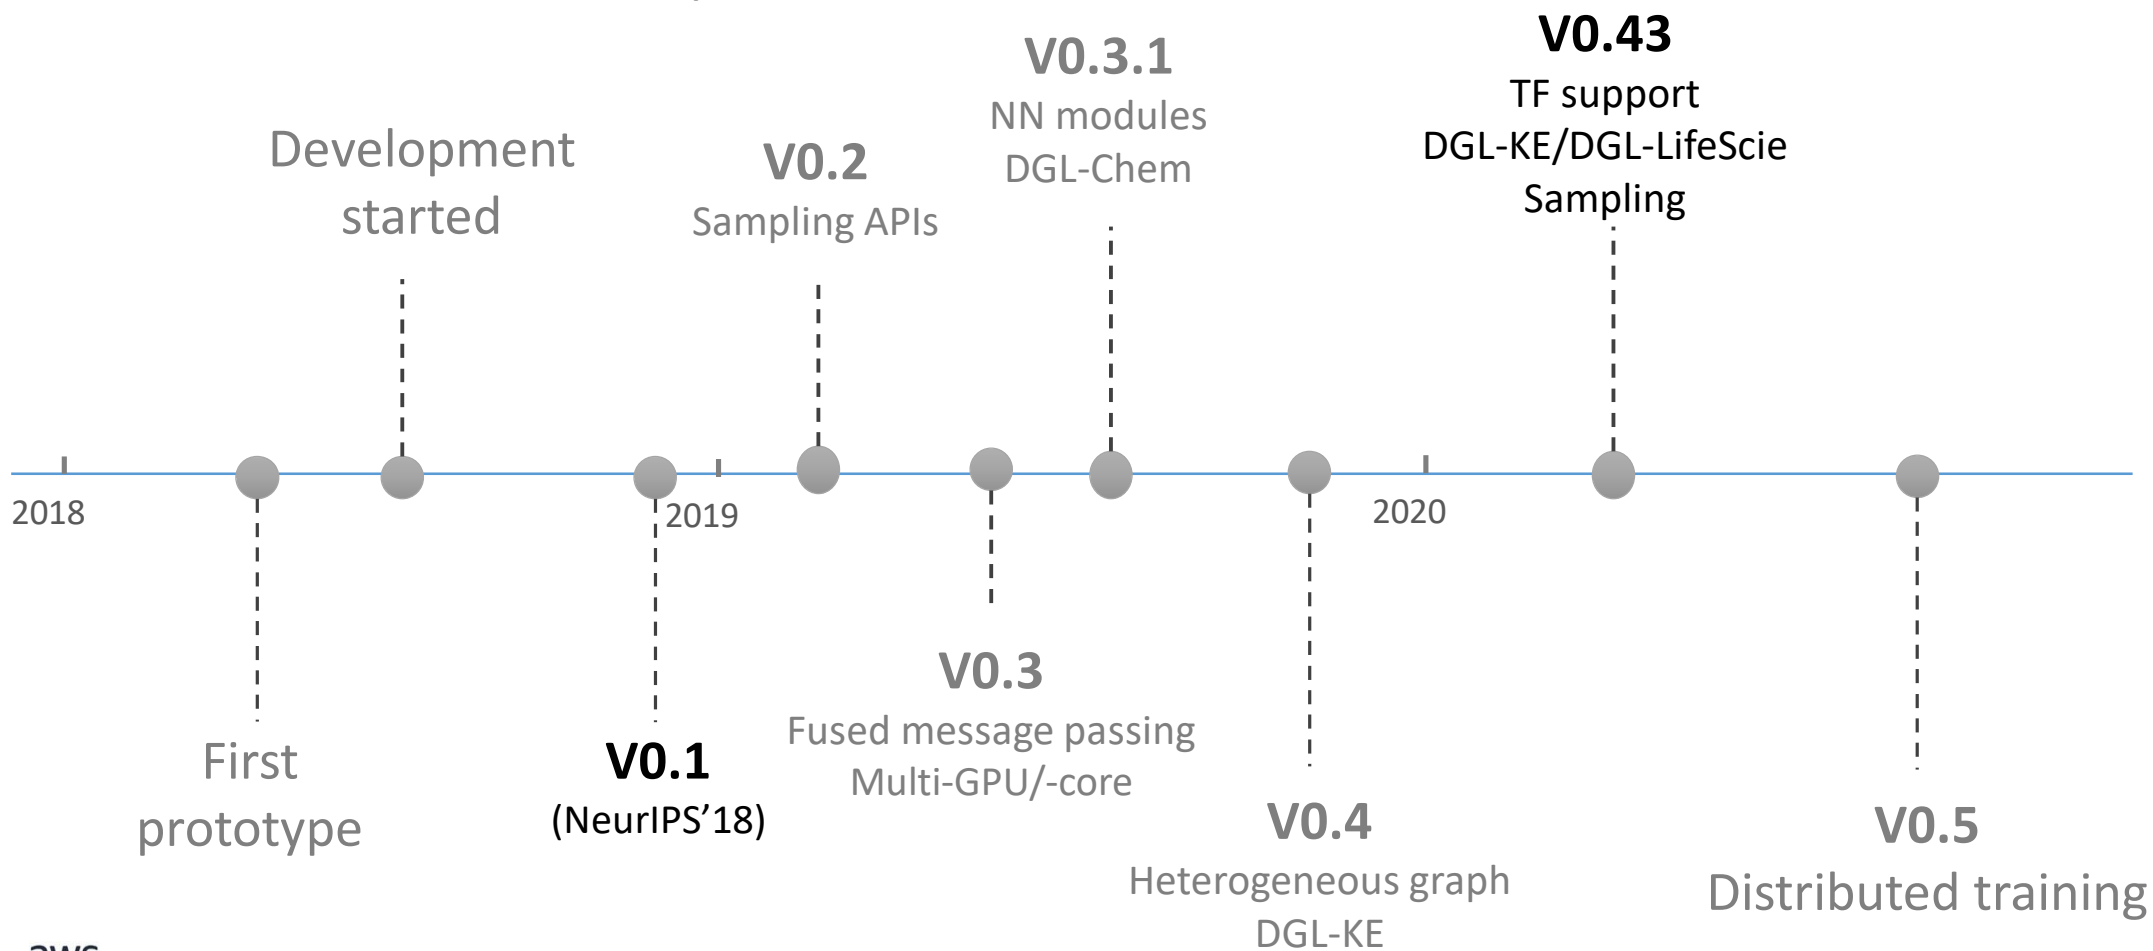


Neurips'19  
WWW'20

3481 4555 github stars  
312k 1770k downloads for all versions on Pip  
8.8K 13K downloads for all version on Conda  
1.8K 6.4K anaconda downloads of 0.4+

32 model examples, 28 NN modules (including  
14 15 GNN convolution modules)  
6 10 pretrained models for chemistry  
GCN, generative, KG, RecSys...  
45 53 contributors, 10 core developers

# DGL: next step(s)



Welcome contributions! Please cite:

## DEEP GRAPH LIBRARY: TOWARDS EFFICIENT AND SCALABLE DEEP LEARNING ON GRAPHS

**Minjie Wang<sup>1 3</sup>, Lingfan Yu<sup>1</sup>, Da Zheng<sup>3</sup>, Quan Gan<sup>4</sup>, Yu Gai<sup>2</sup>, Zihao Ye<sup>4</sup>,  
Mufei Li<sup>4</sup>, Jinjing Zhou<sup>4</sup>, Qi Huang<sup>2</sup>, Chao Ma<sup>4</sup>, Ziyue Huang<sup>5</sup>, Qipeng Guo<sup>6</sup>,  
Hao Zhang<sup>7</sup>, Haibin Lin<sup>3</sup>, Junbo Zhao<sup>1</sup>, Jinyang Li<sup>1</sup>, Alexander Smola<sup>3</sup>, Zheng Zhang<sup>4 2</sup>**

<sup>1</sup> New York University, <sup>2</sup> NYU Shanghai, <sup>3</sup> Amazon Web Services, <sup>4</sup> AWS Shanghai AI Lab

<sup>5</sup> Hong Kong University of Science and Technology, <sup>6</sup> Fudan University,


<sup>7</sup> Chongqing University of Posts and Telecommunications

**George Karypis (Univ Minnesota/AWS)**

We are hiring!

# Q&A

# Discover



# Leverage